

$$1a. H(y) = (0.6 * \log_2(5/3)) + (0.4 * \log_2(5/2)) = 0.971$$

1b.

$$H(x_1) = (0.6 * (0.971 - (0.5 * \log_2(2) + 0.5 * \log_2(2)))) + (0.4 * (0.971 - (0.75 * \log_2(4/3) + 0.25 * \log_2(4))) = 0.6 + (0.4 * 0.811) = 0.06388875022 - 0.0174 = 0.04648875022$$

$$H(x_2) = (0.5 * (0.971 - (0.2 * \log_2(5) + 0.8 * \log_2(5/4)))) + (0.5 * (0.971 - (1 * \log_2(1)))) = (0.5 * 0.722) + (0) = 0.1245359526 + 0.4855 = 0.6100359526$$

$$H(x_3) = (0.3 * (0.971 - (0.33 * \log_2(3) + 0.66 * \log_2(3/2)))) + (0.7 * (0.971 - (4/7 * \log_2(7/4) + 3/7 * \log_2(7/3)))) = 0.01856613729 + -0.009959695224 = 0.008606442066$$

$$H(x_4) = (0.3 * (0.971 - (0.33 * \log_2(3) + 0.66 * \log_2(3/2)))) + (0.7 * (0.971 - (5/7 * \log_2(7/5) + 2/7 * \log_2(7/2)))) = 0.01856613729 + 0.075515602 = 0.09408173929$$

$$H(x_5) = (0.7 * (0.971 - (4/7 * \log_2(7/4) + 3/7 * \log_2(7/3)))) + (0.3 * (0.971 - (0.33 * \log_2(3) + 0.66 * \log_2(3/2)))) = -0.009959695224 + 0.01856613729 = 0.008606442066$$

Splitting on x2 gives the highest information gain.

2a.

```

In [66]: import numpy as np
import mltools as ml
np.random.seed(0)
X = np.genfromtxt('data/X_train.txt', delimiter=',')
Y = np.genfromtxt('data/Y_train.txt', delimiter=',')
X,Y = ml.shuffleData(X,Y)
X = X[:, :41]
print("Minimum of first feature:", min(X[:,0]))
print("Maximum of first feature:", max(X[:,0]))
print("Mean of first feature:", np.mean(X[:,0]))
print("Variance of first feature:", np.var(X[:,0]))

print()
print("Minimum of second feature:", min(X[:,1]))
print("Maximum of second feature:", max(X[:,1]))
print("Mean of second feature:", np.mean(X[:,1]))
print("Variance of second feature:", np.var(X[:,1]))

print()
print("Minimum of third feature:", min(X[:,2]))
print("Maximum of third feature:", max(X[:,2]))
print("Mean of third feature:", np.mean(X[:,2]))
print("Variance of third feature:", np.var(X[:,2]))

print()
print("Minimum of four feature:", min(X[:,3]))
print("Maximum of four feature:", max(X[:,3]))
print("Mean of four feature:", np.mean(X[:,3]))
print("Variance of four feature:", np.var(X[:,3]))

print()
print("Minimum of fifth feature:", min(X[:,4]))
print("Maximum of fifth feature:", max(X[:,4]))
print("Mean of fifth feature:", np.mean(X[:,4]))
print("Variance of fifth feature:", np.var(X[:,4]))
#Xva =
#learner = ml.dtree.treeClassify(Xtr, Ytr, maxDepth=50)

```

```

Minimum of first feature: 0.0
Maximum of first feature: 110285.0
Mean of first feature: 1321.1174134446987
Variance of first feature: 6747189.595085322

```

```

Minimum of second feature: 0.0
Maximum of second feature: 35.0
Mean of second feature: 6.5916745251246125
Variance of second feature: 34.70690630279573

```

```

Minimum of third feature: 0.0
Maximum of third feature: 51536.0
Mean of third feature: 1152.273237235619
Variance of third feature: 5376518.288798102

```

```

Minimum of four feature: 0.0
Maximum of four feature: 21768.0
Mean of four feature: 234.8262548834703

```

Variance of four feature: 260120.83053297663

Minimum of fifth feature: 0.0

Maximum of fifth feature: 27210.0

Mean of fifth feature: 289.75871211100633

Variance of fifth feature: 406615.8651128233

2b.

```
In [67]: #Xtr = X[:int(len(X)/2),:]
#Ytr = Y[:int(len(Y)/2)]
#Xva = X[int(len(X)/2):,:]
#Yva = Y[int(len(Y)/2):]
Xtr,Xva,Ytr,Yva = ml.splitData(X,Y,0.5)
learner = ml.dtree.treeClassify(Xtr, Ytr, maxDepth=50)
Yhat = learner.predict(Xtr)
Yhat2 = learner.predict(Xva)
print()
print("Training Error:", np.sum(Yhat != Ytr)/len(Ytr))
print("Validation Error:", np.sum(Yhat2 != Yva)/len(Yva))
```

Training Error: 0.0

Validation Error: 0.40878469415251956

2c.

```
In [68]: import matplotlib.pyplot as plt
trainerr = []
validerr = []
maxdepth = []
for i in range(16):
    maxdepth.append(i)
    learner = ml.tree.treeClassify(Xtr, Ytr, maxDepth=i)
    Yhat = learner.predict(Xtr)
    Yhat2 = learner.predict(Xva)
    trainerr.append(np.sum(Yhat != Ytr)/len(Ytr))
    validerr.append(np.sum(Yhat2 != Yva)/len(Yva))
    print()
    print("Training Error for max depth", i, ":", np.sum(Yhat != Ytr)/len(Ytr))
    print("Validation Error for max depth", i, ":", np.sum(Yhat2 != Yva)/len(Yva))
plt.plot(maxdepth,trainerr)
plt.plot(maxdepth,validerr)
```

Training Error for max depth 0 : 0.4983836206896552  
 Validation Error for max depth 0 : 0.5125303152789006

Training Error for max depth 1 : 0.39197198275862066  
 Validation Error for max depth 1 : 0.3880355699272433

Training Error for max depth 2 : 0.39197198275862066  
 Validation Error for max depth 2 : 0.3880355699272433

Training Error for max depth 3 : 0.3884698275862069  
 Validation Error for max depth 3 : 0.39099973053085424

Training Error for max depth 4 : 0.36018318965517243  
 Validation Error for max depth 4 : 0.39315548369711667

Training Error for max depth 5 : 0.3464439655172414  
 Validation Error for max depth 5 : 0.3869576933441121

Training Error for max depth 6 : 0.32570043103448276  
 Validation Error for max depth 6 : 0.37510105092966856

Training Error for max depth 7 : 0.30495689655172414  
 Validation Error for max depth 7 : 0.370250606305578

Training Error for max depth 8 : 0.2874461206896552  
 Validation Error for max depth 8 : 0.37294529776340607

Training Error for max depth 9 : 0.2723599137931034  
 Validation Error for max depth 9 : 0.3742926434923201

Training Error for max depth 10 : 0.2572737068965517  
 Validation Error for max depth 10 : 0.37779574238749664

Training Error for max depth 11 : 0.23760775862068967  
 Validation Error for max depth 11 : 0.37752627324171384


Training Error for max depth 12 : 0.21875  
 Validation Error for max depth 12 : 0.3837240635947184

Training Error for max depth 13 : 0.2023168103448276  
Validation Error for max depth 13 : 0.37752627324171384

Training Error for max depth 14 : 0.1869612068965517  
Validation Error for max depth 14 : 0.3769873349501482

Training Error for max depth 15 : 0.17456896551724138  
Validation Error for max depth 15 : 0.3858798167609809

Out[68]: [



Models with higher max depth have higher complexity. Max depth 7 provides the best model.

2d.

```
In [69]: import matplotlib.pyplot as plt
trainerr = []
validerr = []
maxdepth = []
for i in range(14):
    x = pow(2,i)
    maxdepth.append(x)
    learner = ml.tree.treeClassify(Xtr, Ytr, maxDepth=50, minParent = x)
    Yhat = learner.predict(Xtr)
    Yhat2 = learner.predict(Xva)
    trainerr.append(np.sum(Yhat != Ytr)/len(Ytr))
    validerr.append(np.sum(Yhat2 != Yva)/len(Yva))
    print()
    print("Training Error for min parent", x, ":", np.sum(Yhat != Ytr)/len(Ytr))
    print("Validation Error for min parent", x, ":", np.sum(Yhat2 != Yva)/len(Yva))
plt.plot(maxdepth,trainerr)
plt.plot(maxdepth,validerr)
```

Training Error for min parent 1 : 0.0  
 Validation Error for min parent 1 : 0.4125572621934788

Training Error for min parent 2 : 0.0  
 Validation Error for min parent 2 : 0.3971975208838588

Training Error for min parent 4 : 0.009428879310344827  
 Validation Error for min parent 4 : 0.40905416329830235

Training Error for min parent 8 : 0.03879310344827586  
 Validation Error for min parent 8 : 0.4033953112368634

Training Error for min parent 16 : 0.09617456896551724  
 Validation Error for min parent 16 : 0.4074373484236055

Training Error for min parent 32 : 0.15975215517241378  
 Validation Error for min parent 32 : 0.40662894098625707

Training Error for min parent 64 : 0.2200969827586207  
 Validation Error for min parent 64 : 0.39611964430072755

Training Error for min parent 128 : 0.26643318965517243  
 Validation Error for min parent 128 : 0.3856103476151981

Training Error for min parent 256 : 0.30145474137931033  
 Validation Error for min parent 256 : 0.38884397736459175

Training Error for min parent 512 : 0.32112068965517243  
 Validation Error for min parent 512 : 0.36189706278631095


Training Error for min parent 1024 : 0.3407866379310345  
 Validation Error for min parent 1024 : 0.35408245755860956

Training Error for min parent 2048 : 0.36907327586206895  
 Validation Error for min parent 2048 : 0.38426300188628404

Training Error for min parent 4096 : 0.4983836206896552  
 Validation Error for min parent 4096 : 0.5125303152789006

```
Training Error for min parent 8192 : 0.4983836206896552  
Validation Error for min parent 8192 : 0.5125303152789006
```

```
Out[69]: [matplotlib.lines.Line2D at 0x122006340]
```



Models with higher minParent have lower complexity. MinParent 1024 gives the best decision tree model.

2f.

```
In [70]: from sklearn.metrics import roc_curve, auc
import mltools as ml

learner = ml.dtree.treeClassify(Xtr, Ytr, maxDepth=7)

print("Max depth 7 ROC curve")

fpr, tpr, tnr = learner.roc(Xtr, Ytr)
plt.plot(fpr, tpr)

fpr, tpr, tnr = learner.roc(Xva, Yva)
plt.plot(fpr, tpr)


print("Max depth 7 training error AUC:", learner.auc(Xtr, Ytr))
print("Max depth 7 validation error AUC:", learner.auc(Xva, Yva))
print()
print("Max depth 7 ROC curves for training and validation error:")
```

Max depth 5 ROC curve

Max depth 7 training error AUC: 0.7754431445408889

Max depth 7 validation error AUC: 0.6790255696630761

Max depth 7 ROC curves for training and validation error:





```
In [71]: from sklearn.metrics import roc_curve, auc
import mltools as ml

learner = ml.dtree.treeClassify(Xtr, Ytr, maxDepth=50, minParent = 1024)

fpr, tpr, tnr = learner.roc(Xtr,Ytr)
plt.plot(fpr, tpr)


fpr, tpr, tnr = learner.roc(Xva,Yva)
plt.plot(fpr, tpr)
#roc_curve(Yhat, Ytr)

print("Max depth 50 and min parent 1024 training error AUC:", learner.auc(X
print("Max depth 50 and min parent 1024 validation error AUC:", learner.auc
print()
print("Max depth 50 and min parent 1024 ROC curves for training and validat
```

Max depth 50 and min parent 1024 training error AUC: 0.7184301971144076

Max depth 50 and min parent 1024 validation error AUC: 0.69812637943592

Max depth 50 and min parent 1024 ROC curves for training and validation e  
rror:



2g. Kaggle username: Joshua Hsin, AUC: 0.72468

```
In [129]: import numpy as np
import mltools as ml
#X = X[:, :41]

learner = ml.dtree.treeClassify(X, Y, maxDepth=7, minParent = 300)
Xte = np.genfromtxt('data/X_test.txt', delimiter=',')
Yte = np.vstack((np.arange(Xte.shape[0]), learner.predictSoft(Xte)[: ,1])).T

# Output a file with two columns, a row ID and a confidence in class 1:
np.savetxt('Y_submit.txt', Yte, '%d, %.2f', header='Id,Predicted', comments='',
```

```
In [117]: X = np.genfromtxt('data/X_train.txt', delimiter=',')
Y = np.genfromtxt('data/Y_train.txt', delimiter=',')
X,Y = ml.shuffleData(X,Y)
X = X[:, :41]

learner = ml.dtree.treeClassify(X, Y, maxDepth = 7, minParent = 300)

print("AUC:", learner.auc(X,Y))
Yhat = learner.predict(X)
print("Error for optimal classifier:", np.sum(Yhat != Y)/len(Y))
```

AUC: 0.7329601094616423

Error for optimal classifier: 0.3443351744577664

3a. Kaggle username: Joshua Hsin, AUC: 0.70632

```
In [118]: values = [0] * 8
index = [1,5,10,20,25,30,40,50]
values2 = [0] * 8
import random
import numpy as np
#np.random.seed(4)
m,n = Xtr.shape
classifiers = [None] * 50
for i in range(50):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/50))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 50))
for i in range(50):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))
values[7] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 50))
for i in range(50):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[7] = np.sum(predict2 != Ytr)/len(Ytr)
```

Validation Error: 0.39153866882241983

Training Error: 0.3661099137931034

```

In [119]: import random
import numpy as np
#np.random.seed(4)
m,n = Xtr.shape
classifiers = [None] * 40
for i in range(40):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/40))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 40))
for i in range(40):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[6] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 40))
for i in range(40):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[6] = np.sum(predict2 != Ytr)/len(Ytr)

Validation Error: 0.3947722985718135
Training Error: 0.35614224137931033

```

```
In [120]: import random
import numpy as np
#np.random.seed(4)
m,n = Xtr.shape
classifiers = [None] * 30
for i in range(30):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/30))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 30))
for i in range(30):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[5] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 30))
for i in range(30):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[5] = np.sum(predict2 != Ytr)/len(Ytr)

Validation Error: 0.40231743465373215
Training Error: 0.3415948275862069
```

```

In [121]: import random
import numpy as np
#np.random.seed(4)
m,n = Xtr.shape
classifiers = [None] * 25
for i in range(25):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/25))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 25))
for i in range(25):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[4] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 25))
for i in range(25):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[4] = np.sum(predict2 != Ytr)/len(Ytr)

Validation Error: 0.375909458367017
Training Error: 0.34509698275862066

```

```
In [122]: import random
import numpy as np
#np.random.seed(4)
m,n = Xtr.shape
classifiers = [None] * 20
for i in range(20):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/20))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 20))
for i in range(20):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[3] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 20))
for i in range(20):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[3] = np.sum(predict2 != Ytr)/len(Ytr)

Validation Error: 0.3990838049043385
Training Error: 0.34509698275862066
```

```
In [123]: import random
import numpy as np
m,n = Xtr.shape
classifiers = [None] * 10
for i in range(10):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/10))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 10))
for i in range(10):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[2] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 10))
for i in range(10):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[2] = np.sum(predict2 != Ytr)/len(Ytr)
```

Validation Error: 0.38938291565615735

Training Error: 0.33836206896551724

```
In [124]: import random
import numpy as np
m,n = Xtr.shape
classifiers = [None] * 5
for i in range(5):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)/5))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 5))
for i in range(5):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[1] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 5))
for i in range(5):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[1] = np.sum(predict2 != Ytr)/len(Ytr)
```

Validation Error: 0.3858798167609809

Training Error: 0.28502155172413796



```

In [125]: import random
import numpy as np
m,n = Xtr.shape
classifiers = [None] * 1
for i in range(1):
    X,Y = ml.bootstrapData(Xtr, Ytr, int(len(Xtr)))
    ind = random.randint(1,len(Xtr))
    classifiers[i] = ml.dtree.treeClassify(X, Y, maxDepth = 30, minLeaf = 4

mValid = Xva.shape[0]
predict = np.zeros((mValid, 1))
for i in range(1):
    predict[:,i] = classifiers[i].predict(Xva)
predict2 = np.mean(predict,axis = 1)>0.5

print("Validation Error:", np.sum(predict2 != Yva)/len(Yva))

values[0] = np.sum(predict2 != Yva)/len(Yva)

mTrain = Xtr.shape[0]
predict = np.zeros((mTrain, 1))
for i in range(1):
    predict[:,i] = classifiers[i].predict(Xtr)
predict2 = np.mean(predict,axis = 1)>0.5

print("Training Error:", np.sum(predict2 != Ytr)/len(Ytr))
values2[0] = np.sum(predict2 != Ytr)/len(Ytr)

```

Validation Error: 0.4079762867151711


Training Error: 0.19773706896551724

```

In [126]: import matplotlib.pyplot as plt
plt.plot(index, values)
plt.plot(index, values2)

```

Out[126]: [<matplotlib.lines.Line2D at 0x122419760>]



```

In [ ]: 3a.

```

```

In [132]: import random
import numpy as np
Xte = np.genfromtxt('data/X_test.txt', delimiter=',')
X = np.genfromtxt('data/X_train.txt', delimiter=',')
Y = np.genfromtxt('data/Y_train.txt', delimiter=',')
X,Y = ml.shuffleData(X,Y)
X = X[:, :41]
#np.random.seed(4)
m,n = X.shape
classifiers = [None] * 25
for i in range(25):
    Xa,Ya = ml.bootstrapData(X, Y, int(len(X)/25))
    ind = random.randint(1,len(X))
    classifiers[i] = ml.dtree.treeClassify(Xa, Ya, maxDepth = 30,minLeaf =

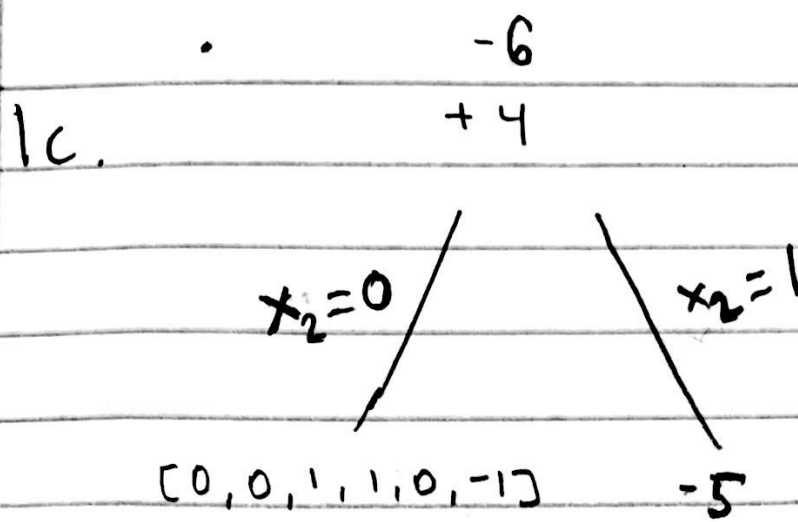
mValid = Xte.shape[0]
predict = np.zeros((mValid, 25))
for i in range(25):
    predict[:,i] = classifiers[i].predict(Xte)
predict2 = np.mean(predict,axis = 1)

Yte = np.vstack((np.arange(Xte.shape[0]), predict2)).T


np.savetxt('Y_submit.txt',Yte,'%d, %.2f',header='Id,Predicted',comments='',

```

Statement of Collaboration: Discussed the difference between entropy and information gain in the context of problem 1.2, discussed the technique for problem 1.3, and discussed logistics for calculating roc and auc in 2f.



[1, 0, 1, 1, 1, +1]  
[0, 0, 1, 0, 0, +1]  
[1, 0, 0, 0, 0, +1]  
[1, 0, 1, 1, 0, +1]



[0, 0, 1, 1, 0, -1]      +3

[0, 0, 1, 0, 0, +1]

+1, -1

